

Data Mining: Learning from Large Data Sets - Fall Semester 2015

member1@student.ethz.ch
member2@student.ethz.ch
member3@student.ethz.ch

October 8, 2015

Approximate near-duplicate search using Locality Sensitive Hashing

Briefly describe the steps used to produce the solution. Feel free to add plots or screenshots if you think it's necessary. The report should contain a maximum of 2 pages.